NLPwDL 2024, Exercise 2

Prof. Dr. Ivan Habernal

2024-10-24

1 Classification evaluation

Accuracy, precision, recall and, F1 measure are typical measures for evaluating the results of machine learning systems. Assume you built a simple multi classification model to solve Part Of Speech tagging which only operates on the three tags `NN`, `VB`, and `ADJ` (noun, verb, and adjective).

Task 1 Compute the classifier's accuracy and precision, recall, and F1 measure for each individual class based on the following confusion matrix:

		predicted class		
		NN	VB	ADJ
	NN	25	5	1
true class	VB	2	15	12
	ADJ	1	6	0

Hint: For n classes and a confusion matrix $C \in \mathbb{R}^{n \times n}$, the evaluation measures are defined for class i by:

$$P_i = \frac{\text{TP}}{\text{TP+FP}} = \frac{C_{i,i}}{\sum_{j=1}^{n} C_{j,i}}$$

$$R_i = \frac{\text{TP}}{\text{TP+FN}} = \frac{C_{i,i}}{\sum_{j=1}^{n} C_{i,j}}$$

and

$$F1 = \frac{2 \cdot P \cdot R}{P + R}$$

$$P_{\text{NN}} = \frac{25}{25 + 3} \qquad\qquad = 0.89$$

$$R_{\text{NN}} = \frac{25}{25 + 6} \qquad\qquad = 0.81$$

$$F1_{\text{NN}} = \frac{2 \cdot 0.89 \cdot 0.81}{0.89 + 0.81} \qquad\qquad = 0.85$$

$$P_{\text{VB}} = \frac{15}{15 + 11} \qquad\qquad = 0.58$$

$$R_{\text{VB}} = \frac{15}{15 + 14} \qquad\qquad = 0.52$$

$$F1_{\text{VB}} = \frac{2 \cdot 0.58 \cdot 0.52}{0.58 + 0.52} \qquad\qquad = 0.55$$

$$P_{\text{ADJ}} = \frac{0}{0 + 13} \qquad\qquad = 0.00$$

$$R_{\text{ADJ}} = \frac{0}{0 + 7} \qquad\qquad = 0.00$$

$$F1_{\text{ADJ}} = \text{undefined}$$

Task 2 Implement a confusion matrix in Python from scratch. You can use `numpy`.

Task 3 Verify your previous hand-crafted calculations. Implement macro-F1 score by (a) averaging F1 score for each class, and (b) by first averaging precision and recall over classes and then computing the F1 score.

Task 4 Pretend you have a highly imbalanced test data of 990 `classA` examples and only 10 `classB` examples. You have two systems: `Model-One` classifies everything as `classA` and `Model-Two` throws a coin for each example and with 50% probability classify the example as `classA` (and as `classB` otherwise). Compute all metrics for both systems.

Task 5 Experiment with classification measures implemented in `scikit-learn`.[1] Focus on F1 score and try several options of `average`: `micro`, `macro`. Compare with your implementation.

---

[1] `https://scikit-learn.org/stable/modules/model_evaluation.html#precision-recall-f-measure-metrics`

Task 6 [Optional, try at home] You might also want to look at Torch-Eval, a lightweight evaluation framework well integrated to the PyTorch environment: `https://pytorch.org/torcheval/`

2 Text generation evaluation

Task 1 Play around with BLEU score using HuggingFace: `https://huggingface.co/spaces/evaluate-metric/bleu`

Task 2 Compare the above to ROUGE metric: `https://huggingface.co/spaces/evaluate-metric/rouge`. Note that this implementation is just a wrapper of another library by Google Research: `https://github.com/google-research/google-research/tree/master/rouge`.