

Natural Language Processing with Deep Learning

Lecture 9 – Text classification 5: Introduction to transformers with BERT

Prof. Dr. Ivan Habernal

December 8, 2023

Natural Language Processing Group
Paderborn University

We focus on Trustworthy Human Language Technologies



www.trusthlt.org

Motivation

Problems: Token (word) embeddings – do not model contextual information

We want contextualized token embeddings

RNN processes left-to-right (or right-to-left)

BERT — The “NLP gamechanger”

Best paper award at NAACL
2019

State-of-the-art results on
various NLP tasks

Directly applicable to other
domains and languages

J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova (2019). “**BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding**”. In: *Proceedings of NAACL*. Minneapolis, Minnesota: Association for Computational Linguistics, pp. 4171–4186

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin Ming-Wei Chang Kenton Lee Kristina Toutanova
Google AI Language
{jacobdevlin, mingweichang, kentonl, kristout}@google.com

Abstract

We introduce a new language representation model called **BERT**, which stands for **B**idirectional **E**ncoder **R**epresentations from **T**ransformers. Unlike recent language repre-

There are two existing strategies for applying pre-trained language representations to downstream tasks: *feature-based* and *fine-tuning*. The feature-based approach, such as **ELMo** (Peters et al., 2018a), uses task-specific architectures that

Neural Machine Translation

Neural Machine Translation

Attention “is all you need”

Multi-task learning

BERT

Input and pre-training

Pre-training

Downstream tasks and fine-tuning

Neural machine translation (NMT)

Why machine translation here?

BERT builds upon techniques from MT

What is machine translation?

- Another popular NLP task
- Many large-scale parallel corpora available



Figure 1: MT is a challenging task!

Recurrent networks for neural MT

Traditionally **encoder-decoder** architectures

- One recurrent neural network processes the entire input and generate its dense representation (**encoder**)
- Other recurrent network produces one token at the time conditioned on the previous states and generated tokens (**decoder**)

Neural MT: Typical architectures (up to 2016-2017)

Long short-term memory (LSTM) / GRU networks

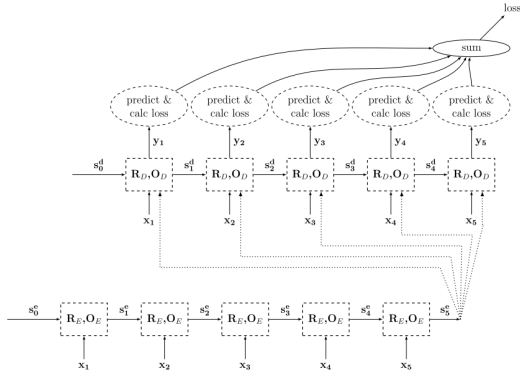


Figure 2: Encoder-decoder RNN

Figure from Y. Goldberg (2016). **"A Primer on Neural Network Models for Natural Language Processing"**. In: *Journal of Artificial Intelligence Research* 57, pp. 345–420

Bottlenecks of RNN for machine translation?

Inherently **sequential** nature

- No parallelization
- Big memory footprint (you must “remember” the entire sequence)
- Long-range dependencies modeling: Distance plays a role!

...but when the goal is to learn a good representation of the input sequence, why not use...

- Convolutional neural networks?

Convolutional neural nets (CNN)

One particular property of CNNs

- Modeling dependencies for a **local context**, but by **stacking layers**, one exactly controls the context size




Figure 3: Receptive field of units in deeper layers is larger

Figure from I. Goodfellow, Y. Bengio, and A. Courville (2016). **Deep Learning**. MIT Press

Convolutional neural nets for MT

CNNs competitive with RNNs for MT¹

- Input tokens as word embeddings (not new) or sub-words (will be explained later)
- Fixed-length input? Set-up a maximum length and use <PAD>ding
- But positional information of tokens is lost...

J. Gehring, M. Auli, D. Grangier, D. Yarats, and Y. N. Dauphin (2017). **“Convolutional Sequence to Sequence Learning”**. In: *Proceedings of the 34th International Conference on Machine Learning*. Ed. by D. Precup and Y. W. Teh. Sydney, Australia: PMLR, pp. 1243–1252

¹Gehring, Auli, Grangier, Yarats, and Dauphin (2017), Facebook AI Research

Convolutional neural nets for MT by Gehring, Auli, Grangier, Yarats, and Dauphin (2017)

Solution: Positional embeddings

- For each input position n , train another embedding vector P_n : $P_1 = (1.12, -78.6, \dots), P_2, \dots, P_N$
- Word embeddings and position embeddings are simply summed up for each input token
- Why? The model knows with which part of the input/output is dealing with
 - Notice: Removing positional embeddings \rightarrow only slightly worse performance

State-of-the-art results and **9.3–21.3 \times faster** than LSTMs on GPU

Attention “is all you need”

Neural Machine Translation

Attention “is all you need”

Multi-task learning

BERT

Input and pre-training

Pre-training

Downstream tasks and fine-tuning

Attention: Modeling dependencies

Recap: How to model long-range dependencies in input?


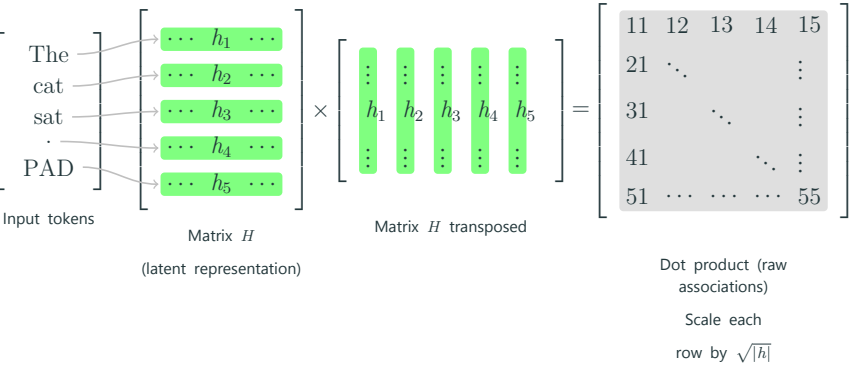


Figure 1: An example of the label consistency problem. Here we would like our model to encourage entities *Albert Einstein* and *Einstein* to get the same label, so as to improve the chance that both are labeled *PERSON*.


- RNNs or stacking CNNs
- **Self-Attention**: Utilize associations between all input word pairs

Figure source: V. Krishnan and C. D. Manning (2006). “**An Effective Two-Stage Model for Exploiting Non-Local Dependencies in Named Entity Recognition**”. In: *Proceedings of ACL*. Sydney, Australia: Association for Computational Linguistics, pp. 1121–1128

Self-Attention in detail (1)




Self-Attention in detail (2)




- Each row corresponds to an input token
- Each row sums up to 1
- Each cell shows the "association strength" with all other tokens

Self-Attention in detail (3)




Each position in the latent representation of a token is weighted by the association strength with other tokens

Self-Attention in detail (4)



Self-Attention in detail: Head



Further operations

- Layer normalization
- Feed-forward layer with ReLU
- Another residual connection and layer normalization

Self-attention: More subtleties

- Run N attention “heads” in parallel and concatenate
- Stack on top of each other M -times

Why self-attention?

- Self-attention layer connects all positions with a constant number of sequentially executed operations
- Recurrent layer requires $O(n)$ sequential operations
- Self-attention layers are **fast**

A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin (2017). **“Attention Is All You Need”**. In: *Advances in Neural Information Processing Systems 30*. Long Beach, CA, USA: Curran Associates, Inc., pp. 5998–6008

Multi-task learning

Neural Machine Translation

Attention “is all you need”

Multi-task learning

BERT

- Input and pre-training


- Pre-training

- Downstream tasks and fine-tuning

Multi-task Learning

Approach to inductive transfer that improves generalization

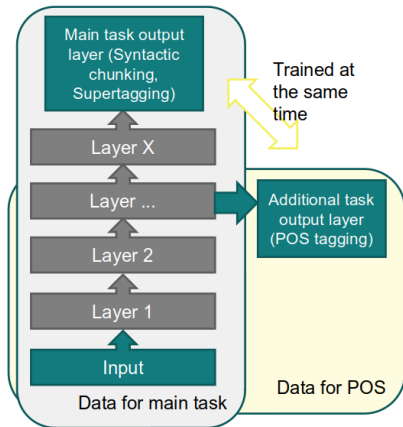
By learning tasks in parallel while using a shared representation



R. Caruana (1997). “**Multi-task Learning**”. In: *Machine Learning* 28.1, pp. 41–75


Multi-task learning in NLP

"In case we suspect the existence of a hierarchy between the different tasks, we show that it is worth-while to incorporate this knowledge in the MTL architecture's design, by making lower level tasks affect the lower levels of the representation."



A. Søgaard and Y. Goldberg (2016). **“Deep multi-task learning with low level tasks supervised at lower layers”.**

Learn a sentence representation on a different task



*"Models learned on NLI can perform better than models trained in unsupervised conditions or on other supervised tasks."*²

²A. Conneau, D. Kiela, H. Schwenk, L. Barrault, and A. Bordes (2017). **"Supervised Learning of Universal Sentence Representations from Natural Language Inference Data"**.

In: *Proceedings of EMNLP*. Copenhagen, Denmark, pp. 670–680

BERT

Neural Machine Translation

Attention “is all you need”

Multi-task learning

BERT


- Input and pre-training

- Pre-training

- Downstream tasks and fine-tuning


BERT: Very abstract view

Input text: *Lorem ipsum dolor*



BERT: The transformer encoder

- Multiple parallel attention "heads" (16 heads)
- With residual connections
- With layer normalization
- Stacked on top of each other (24-times)
- 310,000,000 trainable parameters



Some details (Notation)

Simplify the set notation

$\{1, 2, \dots, N\}$ is a set of integers $1, 2, \dots, N - 1, N$

simplify to $[N]$

For example $t \in [N] \equiv t \in \{1, 2, \dots, N\}$

Notation and formal description of algorithms adopted from M. Phuong and M. Hutter (2022). ***Formal Algorithms for Transformers.*** arXiv: 2207.09238

Note that they use column-vector notation while here (and in all lectures) we use row-vector notation.

Basic single-query attention

Input: $e \in \mathbb{R}^{d_{\text{in}}}$, vector representation of the current token

Input: $e_t \in \mathbb{R}^{d_{\text{in}}}$, vector representations of the context tokens $t \in [T]$

Output: $\tilde{v} \in \mathbb{R}^{d_{\text{out}}}$, vector representation of the token and context combined

Params: $W_q, W_k \in \mathbb{R}^{d_{\text{in}} \times d_{\text{attn}}}$, $b_q, b_k \in \mathbb{R}^{d_{\text{attn}}}$, the query and key linear projections
 $W_v \in \mathbb{R}^{d_{\text{in}} \times d_{\text{out}}}$, $b_v \in \mathbb{R}^{d_{\text{out}}}$, the value linear projection

1: **function** BASIC SINGLE-QUERY ATTENTION

2: $q \leftarrow eW_q + b_q$ ▷ Query linear projection

3: **for** $t \in [T]$ **do**

4: $k_t \leftarrow e_t W_k + b_k$ ▷ Key linear projection

5: $\alpha_t = \frac{\exp(q \cdot k_t / \sqrt{d_{\text{attn}}})}{\sum_{u=1}^T \exp(q \cdot k_u / \sqrt{d_{\text{attn}}})}$ ▷ Softmax over scaled dot products ($\alpha_t \in \mathbb{R}$)

6: $v_t \leftarrow e_t W_v + b_v$ ▷ Value linear projection

7: **return** $\tilde{v} = \sum_{t=1}^T \alpha_t v_t$

Some details

Concatenate matrices of the same dimensions along rows

$$\mathbf{Y} = [\mathbf{X}^1; \mathbf{X}^2; \dots; \mathbf{X}^H] \quad \mathbf{X}^i \in \mathbb{R}^{m \times n} \quad \mathbf{Y} \in \mathbb{R}^{m \times H \cdot n}$$

Example

$$\mathbf{A} = \begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix}, \mathbf{B} = \begin{pmatrix} 11 & 12 \\ 13 & 14 \\ 15 & 16 \end{pmatrix} \quad \mathbf{Y} = [\mathbf{A}; \mathbf{B}] = \begin{pmatrix} 1 & 2 & 11 & 12 \\ 3 & 4 & 13 & 14 \\ 5 & 6 & 15 & 16 \end{pmatrix}$$

Some details

How to add a single vector \mathbf{b} to each row in a matrix \mathbf{W}
($\mathbf{W} \in \mathbb{R}^{m \times n}$, $\mathbf{b} \in \mathbb{R}^n$)

We want $\mathbf{Z} = \mathbf{X} +_{(\text{rows})} \mathbf{b}$

Let $\mathbf{1}^m = (1, 1, \dots, 1_m)$, then $\mathbf{Z} = \mathbf{X} +_{(\text{rows})} \mathbf{b} = \mathbf{X} + (\mathbf{b}^\top \mathbf{1}^m)^\top$

Example

$$\mathbf{X} = \begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix}, \mathbf{b} = (10 \quad 20)$$

$$\mathbf{b}^\top \mathbf{1}^m = \begin{pmatrix} 10 \\ 20 \end{pmatrix} \begin{pmatrix} 1 & 1 & 1 \end{pmatrix} = \begin{pmatrix} 10 & 10 & 10 \\ 20 & 20 & 20 \end{pmatrix} \quad (\mathbf{b}^\top \mathbf{1}^m)^\top = \begin{pmatrix} 10 & 20 \\ 10 & 20 \\ 10 & 20 \end{pmatrix}$$

Some details

Soft-max for matrices row-wise, $\mathbf{A} \in \mathbb{R}^{m \times n}$

$$\text{softmax}_{\text{row}} : \mathbb{R}^{m \times n} \mapsto \mathbb{R}^{m \times n}$$

$$\text{softmax}_{\text{row}}(\mathbf{A})[i, j] = \frac{\exp(\mathbf{A}[i, j])}{\sum_{k=1}^n \exp(\mathbf{A}[i, k])}$$

Bidirectional / unmasked self-attention

Input: $\mathbf{X} \in \mathbb{R}^{\ell_x \times d_x}$, vector representations of the sequence of length ℓ_x

Output: $\tilde{\mathbf{V}} \in \mathbb{R}^{\ell_x \times d_{\text{out}}}$, updated vector representations of tokens in \mathbf{X}

Params \mathcal{W}_{qkv} : $\mathbf{W}_q, \mathbf{W}_k \in \mathbb{R}^{d_x \times d_{\text{attn}}}$, $\mathbf{b}_q, \mathbf{b}_k \in \mathbb{R}^{d_{\text{attn}}}$, $\mathbf{W}_v \in \mathbb{R}^{d_x \times d_{\text{out}}}$, $\mathbf{b}_v \in \mathbb{R}^{d_{\text{out}}}$

1: **function** ATTENTION($\mathbf{X}; \mathcal{W}_{qkv}$)

2: $\mathbf{Q} \leftarrow \mathbf{X} \mathbf{W}_q +_{(\text{rows})} \mathbf{b}_q$

▷ Query $\in \mathbb{R}^{\ell_x \times d_{\text{attn}}}$

3: $\mathbf{K} \leftarrow \mathbf{X} \mathbf{W}_k +_{(\text{rows})} \mathbf{b}_k$

▷ Key $\in \mathbb{R}^{\ell_x \times d_{\text{attn}}}$

4: $\mathbf{V} \leftarrow \mathbf{X} \mathbf{W}_v +_{(\text{rows})} \mathbf{b}_v$

▷ Value $\in \mathbb{R}^{\ell_x \times d_{\text{out}}}$

5: $\mathbf{S} \leftarrow \frac{1}{\sqrt{d_{\text{attn}}}} (\mathbf{Q} \mathbf{K}^\top)$

▷ Scaled score $\in \mathbb{R}^{\ell_x \times \ell_x}$

6: **return** $\tilde{\mathbf{V}} = \text{softmax}_{\text{row}}(\mathbf{S}) \mathbf{V}$

Multi-head bidirectional / unmasked self-attention

Input: $\mathbf{X} \in \mathbb{R}^{\ell_x \times d_x}$, vector representations of the sequence of length ℓ_x

Output: $\tilde{\mathbf{V}} \in \mathbb{R}^{\ell_x \times d_{\text{out}}}$, updated vector representations of tokens in \mathbf{X}

Hyper-param: H , number of attention heads

Params for each $h \in [H]$: \mathcal{W}_{qkv}^h :

- $\mathbf{W}_q^h, \mathbf{W}_k^h \in \mathbb{R}^{d_x \times d_{\text{attn}}}$, $\mathbf{b}_q^h, \mathbf{b}_k^h \in \mathbb{R}^{d_{\text{attn}}}$, $\mathbf{W}_v \in \mathbb{R}^{d_x \times d_{\text{mid}}}$, $\mathbf{b}_v \in \mathbb{R}^{d_{\text{mid}}}$
- $\mathbf{W}_o \in \mathbb{R}^{H \cdot d_{\text{mid}} \times d_{\text{out}}}$, $\mathbf{b}_o \in \mathbb{R}^{d_{\text{out}}}$

1: **function** MHATTENTION($\mathbf{X}; \mathcal{W}$)

2: **for** $h \in [H]$ **do**

3: $\mathbf{Y}^h \leftarrow \text{ATTENTION}(\mathbf{X}; \mathcal{W}_{qkv}^h)$

▷ $\mathbf{Y}^h \in \mathbb{R}^{\ell_x \times d_{\text{mid}}}$

4: $\mathbf{Y} \leftarrow [\mathbf{Y}^1; \mathbf{Y}^2; \dots; \mathbf{Y}^H]$

▷ $\mathbf{Y} \in \mathbb{R}^{\ell_x \times H \cdot d_{\text{mid}}}$

5: **return** $\tilde{\mathbf{V}} = \mathbf{Y}\mathbf{W}_o + \mathbf{b}_o$

Layer normalization

Input: $e \in \mathbb{R}^d$, output of a layer

Input: $\hat{e} \in \mathbb{R}^d$, normalized output of a layer

Parameters: $\gamma, \beta \in \mathbb{R}^d$, element-wise scale and offset

- 1: **function** LAYERNORM($e; \gamma, \beta$)
- 2: $m \leftarrow \frac{1}{d} \sum_{i=1}^d e[i]$ ▷ 'Sample mean' of e
- 3: $v \leftarrow \frac{1}{d} \sum_{i=1}^d (e[i] - m)^2$ ▷ 'Sample variance' of e
- 4: **return** $\hat{e} = \frac{e-m}{\sqrt{v}} \odot \gamma + \beta$ ▷ \odot element-wise product

BERT

Input and pre-training

BERT: Tokenization

Tokenizing into a multilingual WordPiece inventory

- Recall that WordPiece units are sub-word units
- 30,000 WordPiece units (newer models 110k units, 100 languages)

Implications: BERT can "consume" any language

BERT: Input representation

- Each WordPiece token from the input is represented by a **WordPiece embedding** (randomly initialized)
- Each position from the input is associated with a **positional embedding** (also randomly initialized)
- Input length limited to **512** WordPiece tokens, using <PAD>ding
- Special tokens
 - The first token is always a special token **[CLS]**
 - If the task involves two sentences (e.g., NLI), these two sentences are separated by a special token **[SEP]**; also special two **segment position embeddings**

BERT: Input representation summary

Input	[CLS]	my	dog	is	cute	[SEP]	he	likes	play	##ing	[SEP]
Token Embeddings	$E_{[CLS]}$	E_{my}	E_{dog}	E_{is}	E_{cute}	$E_{[SEP]}$	E_{he}	E_{likes}	E_{play}	$E_{\# \# ing}$	$E_{[SEP]}$
	+	+	+	+	+	+	+	+	+	+	+
Segment Embeddings	E_A	E_A	E_A	E_A	E_A	E_A	E_B	E_B	E_B	E_B	E_B
	+	+	+	+	+	+	+	+	+	+	+
Position Embeddings	E_0	E_1	E_2	E_3	E_4	E_5	E_6	E_7	E_8	E_9	E_{10}

BERT

Pre-training

BERT: Self-supervised multi-task pre-training


Prepare two auxiliary tasks that need no labeled data

Task 1: Cloze-test task

- Predict the masked WordPiece unit (multi-class, 30k classes)

Task 2: Consecutive segment prediction

- Did the second text segment appeared after the first segment? (binary)



BERT: Pre-training data generation

Take the entire Wikipedia (in 100 languages; 2,5 billion words)

To generate a single training instance, sample two segments (max combined length 512 WordPiece tokens)

- For Task 2, replace the second segment randomly in 50% (negative samples)
- For Task 1, choose random 15% of the tokens, and in 80% replace with a [MASK]

BERT: Pre-training data – Simplified example

Input = [CLS] the man went to [MASK] store [SEP]

he bought a gallon [MASK] milk [SEP]

Label = IsNext

Input = [CLS] the man [MASK] to the store [SEP]

penguin [MASK] are flight ##less birds [SEP]

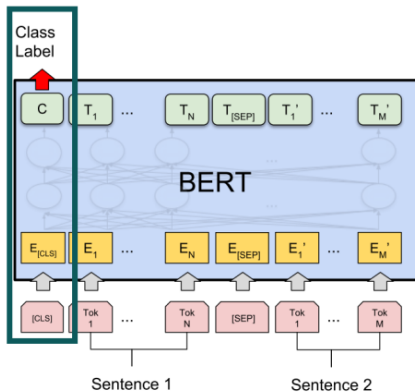
Label = NotNext

- <PAD>ding is missing
- The actual segments are longer and not necessarily actual sentences (just spans)
- The WordPiece tokens match full words / morphology well in this English text, but recall the ones we have seen before

BERT

Downstream tasks and fine-tuning

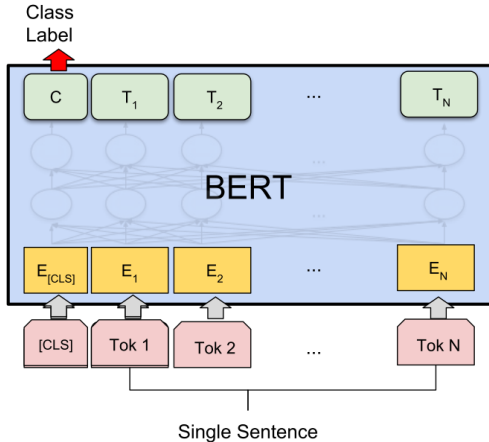
BERT: Representing various NLP tasks



(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG


That explains the special [CLS] token at sequence start

BERT: Representing various NLP tasks



(b) Single Sentence Classification Tasks:
SST-2, CoLA

BERT: Representing various NLP tasks




(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

Not conditioned on surrounding predictions


BERT: Very abstract view

Input text: *Lorem ipsum dolor*



BERT: Very abstract view

Input text: *Lorem ipsum dolor*



P. Izsak, M. Berchansky, and O. Levy (2021). **"How to Train BERT with an Academic Budget"**. In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Online and Punta Cana, Dominican Republic: Association for Computational Linguistics, pp. 10644–10652

Pretraining BERT took originally 4 days on 64 TPUs³

Once pre-trained, transfer and "fine-tune" on your small-data task and get competitive results

³Can be done more efficiently, see, e.g., Izsak, Berchansky, and Levy (2021)

Recap

BERT stays on the shoulders of many clever concepts and techniques, mastered into a single model

What do we know about how BERT works?

“BERTology has clearly come a long way, but it is fair to say we still have more questions than answers about how BERT works.” — Rogers, Kovaleva, and Rumshisky (2020)⁴

A. Rogers, O. Kovaleva, and A. Rumshisky (2020). **“A Primer in BERTology: What We Know About How BERT Works”**.

In: *Transactions of the Association for Computational Linguistics* 8, pp. 842–866

⁴Highly recommended reading!

License and credits

Licensed under Creative Commons
Attribution-ShareAlike 4.0 International
(CC BY-SA 4.0)



Credits

Ivan Habernal

Content from ACL Anthology papers licensed under CC-BY
<https://www.aclweb.org/anthology>