Natural Language Processing with Deep Learning

Lecture 8 — BERT part two

Prof. Dr. Ivan Habernal

December 10, 2025

www.trusthlt.org
Trustworthy Human Language Technologies Group (TrustHLT)
Ruhr University Bochum & Research Center Trustworthy Data Science and Security

RUHR UNIVERSITÄT BOCHUM RUB

TrustHLT 1001100110101

CENTER FOR TRUSTWORTHY DATA SCIENCE AND SECURITY

Where we ended up last week

1 Where we ended up last week
2 Input and pre-training
3 Pre-training
4 Downstream tasks and fine-tuning
5 Sentence BERT
6 What does attention in BERT capture?

BERT: Very abstract view

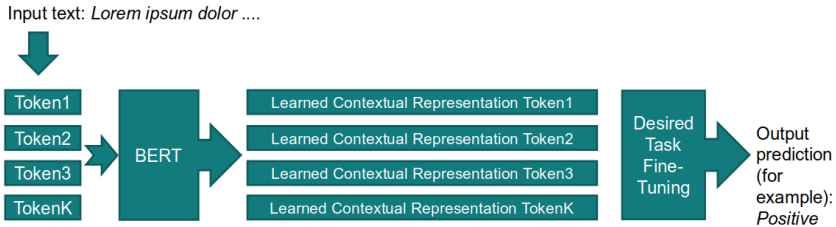

- BERT produces contextualized token embeddings
- BERT can learn them in a ‘clever’ way
- BERT can be applied to many downstream tasks

Input and pre-training

1 Where we ended up last week
2 Input and pre-training
3 Pre-training
4 Downstream tasks and fine-tuning
5 Sentence BERT
6 What does attention in BERT capture?

BERT: Tokenization

Tokenizing into a multilingual WordPiece inventory

- Recall that WordPiece units are sub-word units
- 30,000 WordPiece units (newer models 110k units, 100 languages)

Implications: BERT can "consume" any language

BERT: Input representation

- Each WordPiece token from the input is represented by a **WordPiece embedding** (randomly initialized)
- Each position from the input is associated with a **positional embedding** (also randomly initialized)
- Input length limited to **512** WordPiece tokens, using <PAD>ding
- Special tokens
  - The fist token is always a special token **[CLS]**
  - If the task involves two sentences (e.g., NLI), these two sentences are separated by a special token **[SEP]**; also special two **segment position embeddings**

BERT: Input representation summary

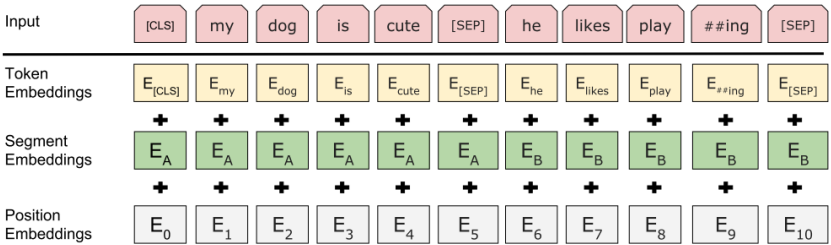

Pre-training

1 Where we ended up last week
2 Input and pre-training
3 Pre-training
4 Downstream tasks and fine-tuning
5 Sentence BERT
6 What does attention in BERT capture?

BERT: Self-supervised multi-task pre-training

Prepare two auxiliary tasks that need no labeled data

Task 1: Cloze-test task

- Predict the masked WordPiece unit (multi-class, 30k classes)

Task 2: Consecutive segment prediction

- Did the second text segment appeared after the first segment? (binary)

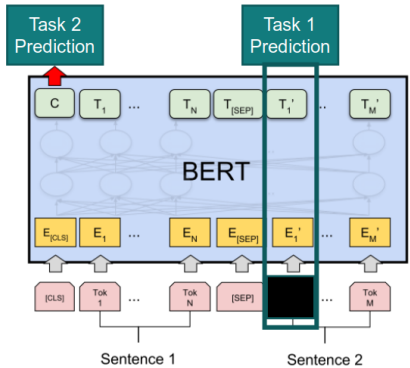

RUHR UNIVERSITÄT BOCHUM RUB

BERT: Pre-training data generation

Take the entire Wikipedia (in 100 languages; 2,5 billion words)

To generate a single training instance, sample two segments (max combined length 512 WordPiece tokens)

- For Task 2, replace the second segment randomly in 50% (negative samples)
- For Task 1, choose random 15% of the tokens, and in 80% replace with a [MASK]

BERT: Pre-training data – Simplified example

Input = [CLS] the man went to [MASK] store [SEP]
he bought a gallon [MASK] milk [SEP]
Label = IsNext

Input = [CLS] the man [MASK] to the store [SEP]
penguin [MASK] are flight ##less birds [SEP]
Label = NotNext

- <PAD>ding is missing
- The actual segments are longer and not necessarily sentences (just spans)
- The WordPiece tokens match full words here

BERT: pre-training by masked language modeling

function ETraining($\{\boldsymbol{x}_n\}_{n=1}^{N_{\text{data}}}$ seqs, $\boldsymbol{\theta}$ init. params; $p_{\text{mask}} \in (0,1)$, N_{epochs}, η)
  for $i \in [N_{\text{epochs}}]$ **do**
    for $n \in [N_{\text{data}}]$ **do**
      $\ell \leftarrow \text{length}(\boldsymbol{x}_n)$
      for $t \in [\ell]$ **do**
        $\tilde{\boldsymbol{x}}_n[t] \leftarrow$ <mask_token> with prob. p_{mask}, otherwise $\boldsymbol{x}_n[t]$
      $\tilde{T} \leftarrow \{t \in [\ell] : \tilde{\boldsymbol{x}}_n[t] =$ <mask_token>$\}$ ▷ Indices of masked tokens
      $\boldsymbol{P}_{\boldsymbol{\theta}} \leftarrow \text{ETransformer}(\tilde{\boldsymbol{x}}_n | \boldsymbol{\theta})$
      $\text{loss}_{\boldsymbol{\theta}} \leftarrow -\sum_{t \in \tilde{T}} \log \boldsymbol{P}_{\boldsymbol{\theta}}[t, \boldsymbol{x}_n[t]]$
      $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \eta \cdot \nabla \text{loss}_{\boldsymbol{\theta}}$
  return $\boldsymbol{\theta}$

Simple example explaining lines 6–7 (masking)

$$\begin{pmatrix}\text{The} & \text{cat} & \text{sat}\end{pmatrix} \rightarrow \boldsymbol{x}_n =$$
$$\begin{pmatrix}21 & 11987 & 5438\end{pmatrix} \quad \text{(Indices in } V\text{)}$$

Random masking (index of `<mask_token>` = 50001):

1. For $t = 1$, the random outcome is "mask"
2. For $t = 2$, the random outcome is "keep"
3. For $t = 3$, the random outcome is "mask"

$$\tilde{\boldsymbol{x}}_n = \begin{pmatrix}50001 & 11987 & 50001\end{pmatrix}, \tilde{T} = \{1, 3\}$$

Explaining line 9 (negative log likelihood)

$\begin{pmatrix}\text{The} & \text{cat} & \text{sat}\end{pmatrix} \rightarrow \boldsymbol{x}_n = \begin{pmatrix}21 & 11987 & 5438\end{pmatrix}, \tilde{\boldsymbol{x}}_n = \begin{pmatrix}50001 & 11987 & 50001\end{pmatrix}, \tilde{T} = \{1, 3\}$

$\boldsymbol{P}_{\boldsymbol{\theta}} \leftarrow \text{ETransformer}(\tilde{\boldsymbol{x}}_n | \boldsymbol{\theta})$

$$\boldsymbol{P}_{\boldsymbol{\theta}} = \begin{pmatrix} 0.001 & 0.0007 & \dots & 0.0003 \\ 0.0013 & 0.0065 & \dots & 0.0001 \\ 0.079 & 0.015 & \dots & 0.0001 \end{pmatrix}$$

$\boldsymbol{P}_{\boldsymbol{\theta}} \in (0, 1)^{\ell_{\text{x}} \times N_{\text{V}}}$, where each row of $\boldsymbol{P}$ is a distribution over the vocabulary

Explaining line 9 (negative log likelihood), $t = 1$

$\boldsymbol{x}_n = (21, 11987, 5438), \tilde{\boldsymbol{x}}_n = (50001, 11987, 50001), \tilde{T} = \{1, 3\}$

$$\boldsymbol{P}_{\boldsymbol{\theta}} = \begin{pmatrix} 0.001 & \dots & 0.0041_{21} & \dots 0.0003 \\ \vdots & & & \end{pmatrix}$$

For $t = 1$, the model should learn to predict "The" (index 21)

Gold: $\boldsymbol{y} = (0, 0, \dots, 1_{21}, \dots, 0) \in \mathbb{R}^{N_\text{V}}$

Pred: $\hat{\boldsymbol{y}} = \boldsymbol{P}_{\boldsymbol{\theta}}[1, :] = (0.001, \dots, 0.0041_{21}, \dots 0.0003) \in \mathbb{R}^{N_\text{V}}$

Recall: Categorical cross entropy loss

$L(\hat{\boldsymbol{y}}, \boldsymbol{y}) := -\sum_{k=1}^{K} \boldsymbol{y}_{[k]} \log\left(\hat{\boldsymbol{y}}_{[k]}\right)$
$= -1 \cdot \log(\hat{\boldsymbol{y}}[21]) = -\log(\boldsymbol{P}_{\boldsymbol{\theta}}[1, 21])$
$= -\log(\boldsymbol{P}_{\boldsymbol{\theta}}[1, \boldsymbol{x}_n[1]]) = -\log(\boldsymbol{P}_{\boldsymbol{\theta}}[t, \boldsymbol{x}_n[t]])$

Explaining line 9 (negative log likelihood), $t = 3$

$\boldsymbol{x}_n = (21, 11987, 5438), \tilde{\boldsymbol{x}}_n = (50001, 11987, 50001), \tilde{T} = \{1, 3\}$

$\boldsymbol{P_\theta} = \begin{pmatrix} \vdots & \dots & \dots \end{pmatrix}$

For $t = 3$, the model should learn to predict "sat" (id 5438)

Categorical cross entropy loss

$L(\hat{\boldsymbol{y}}, \boldsymbol{y}) := -\sum_{k=1}^{K} \boldsymbol{y}_{[k]} \log\left(\hat{\boldsymbol{y}}_{[k]}\right)$
$= -1 \cdot \log(\hat{\boldsymbol{y}}[5438]) = -\log(\boldsymbol{P_\theta}[3, 5438]) = -\log(\boldsymbol{P_\theta}[t, \boldsymbol{x}_n[t]])$

Sum over all masked token positions in $\tilde{T}$ gives us line 9:

$$\text{loss}_{\boldsymbol{\theta}} \leftarrow -\sum_{t \in \tilde{T}} \log \boldsymbol{P_\theta}[t, \boldsymbol{x}_n[t]]$$

Downstream tasks and fine-tuning

1 Where we ended up last week
2 Input and pre-training
3 Pre-training
4 Downstream tasks and fine-tuning
5 Sentence BERT
6 What does attention in BERT capture?

BERT: Representing various NLP tasks

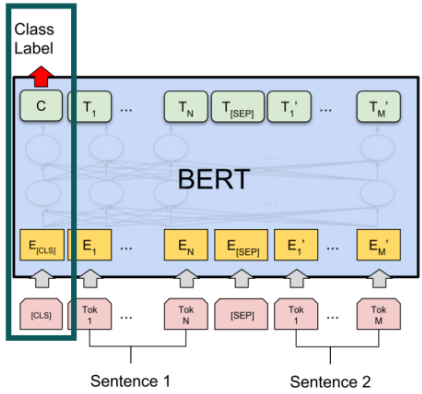

(a) Sentence Pair Classification Tasks: MNLI, QQP, QNLI, STS-B, MRPC, RTE, SWAG

That explains the special [CLS] token at sequence start

BERT: Representing various NLP tasks

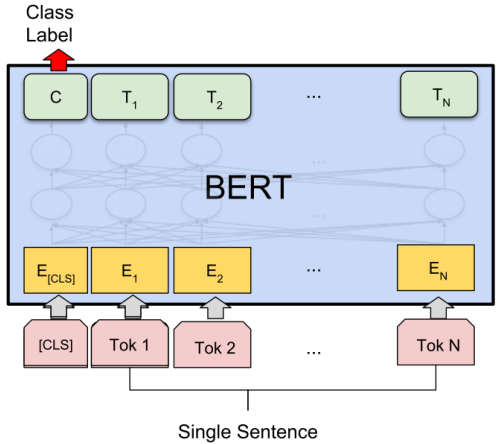

(b) Single Sentence Classification Tasks: SST-2, CoLA

BERT: Representing various NLP tasks

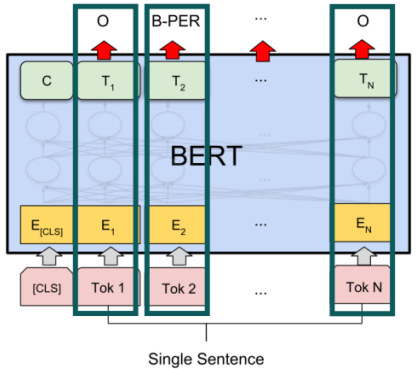

(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

Not conditioned on surrounding predictions

BERT pre-training time

P. Izsak, M. Berchansky, and O. Levy (2021). **"How to Train BERT with an Academic Budget".** In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing.* Online and Punta Cana, Dominican Republic: Association for Computational Linguistics, pp. 10644–10652

Pretraining BERT took originally 4 days on 64 TPUs[1]

Once pre-trained, transfer and "fine-tune" on your small-data task and get competitive results

[1] Can be done more efficiently, see, e.g., Izsak, Berchansky, and Levy (2021)

Sentence BERT

1. Where we ended up last week
2. Input and pre-training
3. Pre-training
4. Downstream tasks and fine-tuning
5. Sentence BERT
6. What does attention in BERT capture?

Motivation for Sentence BERT

BERT — state-of-the-art performance on sentence-pair tasks (e.g, semantic textual similarity), but

- both sentences fed into the network → massive computational overhead
- "finding the most similar pair in a collection of 10,000 sentences requires about 50 million inference computations (65 hours) with BERT"

BERT unsuitable for semantic similarity search or unsupervised tasks like clustering

N. Reimers and I. Gurevych (2019). **"Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks".** In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*. Hong Kong, China: Association for Computational Linguistics, pp. 3980–3990

Sometimes we need just ‘suitable’ sentence embeddings

For **semantic clustering** or **semantic search** — we map each sentence to a vector space

- semantically similar sentences are close to each other

How to use BERT for that?

- average the BERT output layer (known as BERT embeddings)
- use the first token (the [CLS] token) output

→ this practice yields rather bad sentence embeddings, often worse than averaging static word embeddings

N. Reimers and I. Gurevych (2019). **“Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks”.** In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*. Hong Kong, China: Association for Computational Linguistics, pp. 3980–3990

Fine-tune S-BERT on sentence pair classification

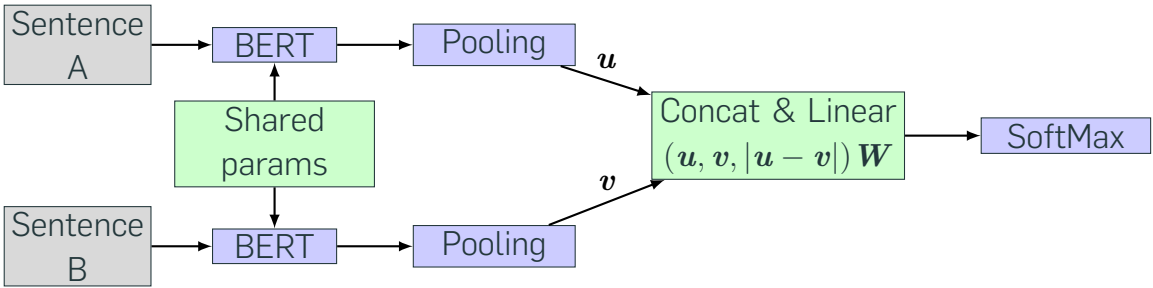

Training data:

- SNLI (570,000 sentence pairs)
- MultiNLI (430,000 sentence pairs)

S-BERT provided superior sentence embeddings

Model	MR	CR	SUBJ	MPQA	SST	TREC	MRPC	Avg.
Avg. GloVe embeddings	77.25	78.30	91.17	87.85	80.18	83.0	72.87	81.52
Avg. fast-text embeddings	77.96	79.23	91.68	87.81	82.15	83.6	74.49	82.42
Avg. BERT embeddings	78.66	86.25	94.37	88.66	84.40	92.8	69.45	84.94
BERT CLS-vector	78.68	84.85	94.21	88.23	84.13	91.4	71.13	84.66
InferSent - GloVe	81.57	86.54	92.50	**90.38**	84.18	88.2	75.77	85.59
Universal Sentence Encoder	80.09	85.19	93.98	86.70	86.38	**93.2**	70.14	85.10
SBERT-NLI-base	83.64	89.43	94.39	89.86	88.96	89.6	**76.00**	87.41
SBERT-NLI-large	**84.88**	**90.07**	**94.52**	90.33	**90.66**	87.4	75.94	**87.69**

N. Reimers and I. Gurevych (2019). **"Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks".** In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*. Hong Kong, China: Association for Computational Linguistics, pp. 3980–3990

Figure 1: Evaluation of SBERT sentence embeddings using the SentEval toolkit. SentEval evaluates sentence embeddings on different sentence classification tasks by training a logistic regression classifier using the sentence embeddings as features.

S-BERT efficiency

"For example, clustering of 10,000 sentences with hierarchical clustering requires with BERT about 65 hours (50 Million sentence combinations). With SBERT, we were able to reduce the effort to about 5 seconds."

N. Reimers and I. Gurevych (2019). **"Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks".** In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*. Hong Kong, China: Association for Computational Linguistics, pp. 3980–3990

What does attention in BERT capture?

1 Where we ended up last week
2 Input and pre-training
3 Pre-training
4 Downstream tasks and fine-tuning
5 Sentence BERT
6 What does attention in BERT capture?

Attention maps to study what linguistic features BERT learns

K. Clark, U. Khandelwal, O. Levy, and C. D. Manning (2019). **"What Does BERT Look at? An Analysis of BERT's Attention".** In: *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*. Florence, Italy: Association for Computational Linguistics, pp. 276–286

Attention

- is naturally interpretable
- attention weight has a clear meaning: how much a particular word will be weighted when computing the next representation for the current word

Explore syntactic capabilities of attention heads

Clark, Khandelwal, Levy, and Manning (2019) probe each attention head for linguistic phenomena

- treat each attention head as a simple no-training-required classifier that, given a word as input, outputs the most attended-to other word
- evaluate the ability of the heads to classify various syntactic relations

K. Clark, U. Khandelwal, O. Levy, and C. D. Manning (2019). **"What Does BERT Look at? An Analysis of BERT's Attention".** In: *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*. Florence, Italy: Association for Computational Linguistics, pp. 276–286

BERT learns to encode syntax quite naturally

"We find heads that find direct objects of verbs, determiners of nouns, objects of prepositions, and objects of possesive pronouns with >75% accuracy."

"These results are **intriguing** *because the behavior of the attention heads* **emerges purely from self-supervised training on unlabeled data**, *without explicit supervision for syntax or coreference."*

K. Clark, U. Khandelwal, O. Levy, and C. D. Manning (2019). **"What Does BERT Look at? An Analysis of BERT's Attention".** In: *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*. Florence, Italy: Association for Computational Linguistics, pp. 276–286

Example: Layer 8, attention head 11

- **Noun modifiers** (e.g., determiners) attend to their noun
- 94.3% accuracy at the `det` relation

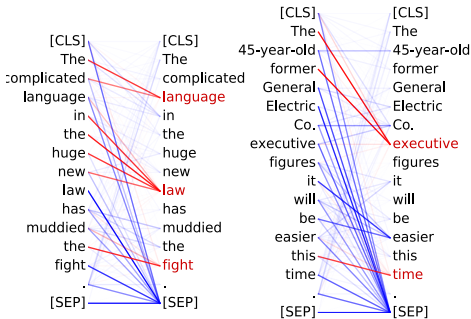

K. Clark, U. Khandelwal, O. Levy, and C. D. Manning (2019). **"What Does BERT Look at? An Analysis of BERT's Attention".** In: *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*. Florence, Italy: Association for Computational Linguistics, pp. 276–286

Figure 2: The darkness of a line indicates the strength of the attention weight. All attention to/from red words is colored red.

Recap

BERT stays on the shoulders of many clever concepts and techniques, mastered into a single model

What do we know about how BERT works?

"BERTology has clearly come a long way, but it is fair to say we still have more questions than answers about how BERT works." — Rogers, Kovaleva, and Rumshisky (2020)[2]

A. Rogers, O. Kovaleva, and A. Rumshisky (2020). **"A Primer in BERTology: What We Know About How BERT Works".** In: *Transactions of the Association for Computational Linguistics* 8, pp. 842–866

[2] Highly recommended reading!

License and credits

Licensed under Creative Commons Attribution-ShareAlike 4.0 International (CC BY-SA 4.0)

Credits

Ivan Habernal

Content from ACL Anthology papers licensed under CC-BY
`https://www.aclweb.org/anthology`